


# HomeWork\_4


Viveksinh Solanki

4/7/2020

```
### Problem 1 ###  
  
## a)  
weekly_df <- read.csv('Weekly.csv')  
#View(weekly_df)  
  
# plots  
plot(weekly_df$Today)
```



```
boxplot(weekly_df$Today)
```



*# From the plots, we can see that value for 'today' variable is lying between  
# -5 and 5. From the scatter plot we can see that straight line can be fitted  
# through whole dataset for 'today' as target variable*

```
# summary stats
summary(weekly_df)
```

```
##      Year      Lag1      Lag2      Lag3
## Min.   :1990  Min.   :-18.1950  Min.   :-18.1950  Min.   :-18.1950
## 1st Qu.:1995  1st Qu.: -1.1540  1st Qu.: -1.1540  1st Qu.: -1.1580
## Median :2000  Median :  0.2410  Median :  0.2410  Median :  0.2410
## Mean   :2000  Mean   :  0.1506  Mean   :  0.1511  Mean   :  0.1472
## 3rd Qu.:2005  3rd Qu.:  1.4050  3rd Qu.:  1.4090  3rd Qu.:  1.4090
## Max.    :2010  Max.    : 12.0260  Max.    : 12.0260  Max.    : 12.0260
##      Lag4      Lag5      Volume
## Min.   :-18.1950  Min.   :-18.1950  Min.   :0.08747
## 1st Qu.: -1.1580  1st Qu.: -1.1660  1st Qu.:0.33202
## Median :  0.2380  Median :  0.2340  Median :1.00268
## Mean   :  0.1458  Mean   :  0.1399  Mean   :1.57462
## 3rd Qu.:  1.4090  3rd Qu.:  1.4050  3rd Qu.:2.05373
## Max.    : 12.0260  Max.    : 12.0260  Max.    :9.32821
##      Today      Direction
## Min.   :-18.1950  Down:484
## 1st Qu.: -1.1540  Up :605
## Median :  0.2410
## Mean   :  0.1499
```

```

## 3rd Qu.: 1.4050
## Max.    : 12.0260

# From the summary stats, it is visible that except volume all other features
# have mean value ~0.15
# number of data points with 'Down' direction is 484
# number of data points with 'Up' direction is 605

## b) Logistic Regression
library(ISLR)

glm.fit=glm(weekly_df$Direction~Lag1+Lag2+Lag3+Lag4+Lag5+Volume,
            data=weekly_df,family=binomial)

summary(glm.fit)

##
## Call:
## glm(formula = weekly_df$Direction ~ Lag1 + Lag2 + Lag3 + Lag4 +
##      Lag5 + Volume, family = binomial, data = weekly_df)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6949  -1.2565   0.9913   1.0849   1.4579
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.26686    0.08593   3.106  0.0019 **
## Lag1        -0.04127    0.02641  -1.563  0.1181
## Lag2         0.05844    0.02686   2.175  0.0296 *
## Lag3        -0.01606    0.02666  -0.602  0.5469
## Lag4        -0.02779    0.02646  -1.050  0.2937
## Lag5        -0.01447    0.02638  -0.549  0.5833
## Volume      -0.02274    0.03690  -0.616  0.5377
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1496.2  on 1088  degrees of freedom
## Residual deviance: 1486.4  on 1082  degrees of freedom
## AIC: 1500.4
##
## Number of Fisher Scoring iterations: 4

confint(glm.fit)

## Waiting for profiling to be done...
##
##              2.5 %    97.5 %
## (Intercept)  0.098808746 0.43580101
## Lag1        -0.093477110 0.01029269
## Lag2         0.006197597 0.11169774
## Lag3        -0.068653910 0.03604309
## Lag4        -0.079952378 0.02401603

```

```
## Lag5      -0.066495108 0.03711989
## Volume    -0.095051949 0.04979338
```

```
confint.default(glm.fit)
```

```
##              2.5 %      97.5 %
## (Intercept) 0.098445204 0.43528308
## Lag1        -0.093032105 0.01049422
## Lag2         0.005787254 0.11109610
## Lag3        -0.068319640 0.03619735
## Lag4        -0.079657357 0.02407694
## Lag5        -0.066185275 0.03724115
## Volume      -0.095060526 0.04957746
```

```
exp(coef(glm.fit))
```

```
## (Intercept)      Lag1      Lag2      Lag3      Lag4      Lag5
## 1.3058630 0.9595710 1.0601831 0.9840671 0.9725924 0.9856322
## Volume
## 0.9775151
```

```
# From above results, predictor Lag2 seems statistically significant.
```

```
## c)
```

```
glm.probs=predict(glm.fit, type="response")
glm.probs[1:5]
```

```
##      1      2      3      4      5
## 0.6086249 0.6010314 0.5875699 0.4816416 0.6169013
```

```
dim(weekly_df)
```

```
## [1] 1089 9
```

```
glm.pred=rep("Down",dim(weekly_df)[1])
glm.pred[glm.probs>0.5]="Up"
contrasts(weekly_df$Direction)
```

```
##      Up
## Down 0
## Up   1
```

```
# Confusion matrix
```

```
table(glm.pred,weekly_df$Direction)
```

```
##
## glm.pred Down Up
## Down 54 48
## Up 430 557
```

```
# Overall fraction of correct predictions
```

```
mean(glm.pred == weekly_df$Direction)
```

```
## [1] 0.5610652
```

```
#overall error rate
```

```
1-mean(glm.pred == weekly_df$Direction)
```

```
## [1] 0.4389348
```

```

# Understanding confusion matrix and mistakes made by logistic regression
#error among direction 'Up'
48/(48+557)

## [1] 0.07933884

#sensitivity (percentage of true 'Up' direction identified)
557/(48+557) ## (tp/tp+fn)

## [1] 0.9206612

#specificity (percentage of true 'Down' direction that are correctly identified)
54/(54+430) ## (tn/tn+fp)

## [1] 0.1115702

## d)

# split data into training and held out/testing sets
training=weekly_df[1:985,]
test=weekly_df[986:1089,]

glm.fit=glm(Direction~Lag2,data=training,family=binomial)
summary(glm.fit)

##
## Call:
## glm(formula = Direction ~ Lag2, family = binomial, data = training)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.536  -1.264   1.021   1.091   1.368
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.20326    0.06428   3.162  0.00157 **
## Lag2         0.05810    0.02870   2.024  0.04298 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1354.7  on 984  degrees of freedom
## Residual deviance: 1350.5  on 983  degrees of freedom
## AIC: 1354.5
##
## Number of Fisher Scoring iterations: 4

glm.probs=predict(glm.fit,newdata = test, type="response")
glm.pred=rep("Down",dim(test)[1])
glm.pred[glm.probs>0.5]="Up"
contrasts(test$Direction)

##      Up
## Down  0
## Up    1

```

```

# Confusion matrix
table(glm.pred,test$Direction)

##
## glm.pred Down Up
##      Down    9  5
##      Up     34 56

# Overall fraction of correct predictions
logistic_acc <- mean(glm.pred == test$Direction)
logistic_acc

## [1] 0.625

#overall error rate
1-mean(glm.pred == test$Direction)

## [1] 0.375

## e) LDA
library(MASS)

lda.fit=lda(Direction~Lag2,data=training)
lda.fit

## Call:
## lda(Direction ~ Lag2, data = training)
##
## Prior probabilities of groups:
##      Down      Up
## 0.4477157 0.5522843
##
## Group means:
##      Lag2
## Down -0.03568254
## Up    0.26036581
##
## Coefficients of linear discriminants:
##      LD1
## Lag2 0.4414162

lda.pred=predict(lda.fit,test)
lda.class=lda.pred$class
contrasts(test$Direction)

##      Up
## Down  0
## Up    1

# Confusion matrix
table(lda.class ,test$Direction)

##
## lda.class Down Up
##      Down    9  5
##      Up     34 56

```

```

# Overall fraction of correct predictions
lda_acc <- mean(lda.class == test$Direction)
lda_acc

## [1] 0.625

## f) QDA

qda.fit=qda(Direction~Lag2,data=training)
qda.fit

## Call:
## qda(Direction ~ Lag2, data = training)
##
## Prior probabilities of groups:
##      Down      Up
## 0.4477157 0.5522843
##
## Group means:
##      Lag2
## Down -0.03568254
## Up    0.26036581

qda.pred=predict(qda.fit,test)
qda.class=qda.pred$class
contrasts(test$Direction)

##      Up
## Down  0
## Up    1

# Confusion matrix
table(qda.class ,test$Direction)

##
## qda.class Down Up
##      Down    0  0
##      Up    43 61

# Overall fraction of correct predictions
qda_acc <- mean(qda.class == test$Direction)
qda_acc

## [1] 0.5865385

## g) KNN with k=1
test.x=cbind(test$Lag2)
training.x=cbind(training$Lag2)
library(class)
knn.pred=knn(training.x,test.x,training$Direction,k=1)

contrasts(test$Direction)

##      Up
## Down  0
## Up    1

```

```

# Confusion matrix
table(knn.pred ,test$Direction)

##
## knn.pred Down Up
##      Down   21 30
##      Up     22 31

# Overall fraction of correct predictions
knn_acc <- mean(knn.pred == test$Direction)
knn_acc

## [1] 0.5

## h) Compare above results - Logistic regression and LDA provide best results
# Predictions
comparison=cbind(test,glm.pred,lda.class,qda.class,knn.pred)
head(comparison)

##      Year  Lag1  Lag2  Lag3  Lag4  Lag5  Volume  Today Direction
## 986 2009  6.760 -1.698  0.926  0.418 -2.251  3.793110 -4.448      Down
## 987 2009 -4.448  6.760 -1.698  0.926  0.418  5.043904 -4.518      Down
## 988 2009 -4.518 -4.448  6.760 -1.698  0.926  5.948758 -2.137      Down
## 989 2009 -2.137 -4.518 -4.448  6.760 -1.698  6.129763 -0.730      Down
## 990 2009 -0.730 -2.137 -4.518 -4.448  6.760  5.602004  5.173       Up
## 991 2009  5.173 -0.730 -2.137 -4.518 -4.448  6.217632 -4.808      Down
##      glm.pred lda.class qda.class knn.pred
## 986      Up      Up      Up      Up
## 987      Up      Up      Up      Up
## 988      Down     Down     Up     Down
## 989      Down     Down     Up     Down
## 990      Up      Up      Up     Down
## 991      Up      Up      Up      Up

# Overall fraction of correct predictions
cbind(logistic_acc, lda_acc, qda_acc, knn_acc)

##      logistic_acc lda_acc  qda_acc knn_acc
## [1,]      0.625   0.625 0.5865385      0.5

## i) Experiments with different predictors

#-- A) predictors: Lag2 + Volume
col_names <- c('Lag2', 'Volume', 'Direction')
a_subset <- weekly_df[col_names]
training=a_subset[1:985,]
test=a_subset[986:1089,]

#-- I) Logistic Regression
glm.fit=glm(Direction~Lag2+Volume, data=training,family=binomial)
summary(glm.fit)

##
## Call:
## glm(formula = Direction ~ Lag2 + Volume, family = binomial, data = training)
##
## Deviance Residuals:

```



```

##      Min      1Q  Median      3Q      Max
## -1.413 -1.262  1.020   1.087   1.485
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.27069    0.09006   3.006  0.00265 **
## Lag2         0.05350    0.02905   1.842  0.06554 .
## Volume      -0.05548    0.05180  -1.071  0.28417
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1354.7  on 984  degrees of freedom
## Residual deviance: 1349.4  on 982  degrees of freedom
## AIC: 1355.4
##
## Number of Fisher Scoring iterations: 4

glm.probs=predict(glm.fit,newdata = test, type="response")
glm.pred=rep('Down',dim(test)[1])
glm.pred[glm.probs>0.5]='Up'

# Confusion matrix
table(glm.pred,test$Direction)

##
## glm.pred Down Up
##      Down   20 25
##      Up    23 36

# Overall fraction of correct predictions
logistic_acc <- mean(glm.pred == test$Direction)
logistic_acc

## [1] 0.5384615

#-- II) LDA
lda.fit=lda(Direction~Lag2+Volume,
             data=training)
lda.fit

## Call:
## lda(Direction ~ Lag2 + Volume, data = training)
##
## Prior probabilities of groups:
##      Down      Up
## 0.4477157 0.5522843
##
## Group means:
##      Lag2      Volume
## Down -0.03568254 1.266966
## Up    0.26036581 1.156529
##
## Coefficients of linear discriminants:
##      LD1

```

```

## Lag2    0.3590979
## Volume -0.3756392

lda.pred=predict(lda.fit,test)
lda.class=lda.pred$class
# Confusion matrix
table(lda.class ,test$Direction)

##
## lda.class Down Up
##      Down   20 25
##      Up     23 36
# Overall fraction of correct predictions
lda_acc <- mean(lda.class == test$Direction)
lda_acc

## [1] 0.5384615

#-- III) QDA
qda.fit=qda(Direction~Lag2+Volume,
             data=training)
qda.fit

## Call:
## qda(Direction ~ Lag2 + Volume, data = training)
##
## Prior probabilities of groups:
##      Down      Up
## 0.4477157 0.5522843
##
## Group means:
##      Lag2   Volume
## Down -0.03568254 1.266966
## Up    0.26036581 1.156529
qda.pred=predict(qda.fit,test)
qda.class=qda.pred$class
# Confusion matrix
table(qda.class ,test$Direction)

##
## qda.class Down Up
##      Down   32 44
##      Up     11 17
# Overall fraction of correct predictions
qda_acc <- mean(qda.class == test$Direction)
qda_acc

## [1] 0.4711538

#-- IV) KNN for k=1,2,3,5,7

test.x=cbind(test$Lag2, test$Volume)
training.x=cbind(training$Lag2, training$Volume)
library(class)

```

```

k_ls <- c(1,2,3,5,7)
k_acc <- c()
i=0

for(k in k_ls){
  knn.pred=knn(training.x,test.x,training$Direction,k=k)
  contrasts(test$Direction)
  # Confusion matrix
  table(knn.pred ,test$Direction)
  # Overall fraction of correct predictions
  knn_acc <- mean(knn.pred == test$Direction)
  k_acc <- c(k_acc,knn_acc)
  i=i+1
}

k_acc

## [1] 0.5576923 0.4711538 0.5480769 0.5288462 0.4711538

#-- B) predictors: Lag1 + Lag3 + Lag4 + Lag5 + Volume
col_names <- c('Lag1', 'Lag3', 'Lag4', 'Lag5', 'Volume', 'Direction')
a_subset <- weekly_df[col_names]
training=a_subset[1:985,]
test=a_subset[986:1089,]

#-- I) Logistic Regression
glm.fit=glm(Direction~Lag1+Lag3+Lag4+Lag5+Volume, data=training,family=binomial)
summary(glm.fit)

##
## Call:
## glm(formula = Direction ~ Lag1 + Lag3 + Lag4 + Lag5 + Volume,
##      family = binomial, data = training)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.5522  -1.2490   0.9887   1.0868   1.4014
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.35621    0.09286   3.836 0.000125 ***
## Lag1        -0.06643    0.02925  -2.271 0.023120 *
## Lag3        -0.01972    0.02907  -0.678 0.497565
## Lag4        -0.02975    0.02908  -1.023 0.306292
## Lag5        -0.04051    0.02898  -1.398 0.162104
## Volume      -0.10376    0.05312  -1.953 0.050793 .
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 1354.7  on 984  degrees of freedom
## Residual deviance: 1344.6  on 979  degrees of freedom
## AIC: 1356.6

```

```

##
## Number of Fisher Scoring iterations: 4

glm.probs=predict(glm.fit,newdata = test, type="response")
glm.pred=rep('Down',dim(test)[1])
glm.pred[glm.probs>0.5]='Up'

# Confusion matrix
table(glm.pred,test$Direction)

##
## glm.pred Down Up
##      Down   36 47
##      Up     7 14

# Overall fraction of correct predictions
logistic_acc <- mean(glm.pred == test$Direction)
logistic_acc

## [1] 0.4807692

#-- II) LDA
lda.fit=lda(Direction~Lag1+Lag3+Lag4+Lag5+Volume,
             data=training)
lda.fit

## Call:
## lda(Direction ~ Lag1 + Lag3 + Lag4 + Lag5 + Volume, data = training)
##
## Prior probabilities of groups:
##      Down      Up
## 0.4477157 0.5522843
##
## Group means:
##      Lag1      Lag3      Lag4      Lag5      Volume
## Down  0.28944444 0.17080045 0.15925624 0.21409297 1.266966
## Up    -0.009213235 0.08404044 0.09220956 0.04548897 1.156529
##
## Coefficients of linear discriminants:
##      LD1
## Lag1   -0.32218170
## Lag3   -0.09604607
## Lag4   -0.14492618
## Lag5   -0.19661300
## Volume -0.50985059

lda.pred=predict(lda.fit,test)
lda.class=lda.pred$class
# Confusion matrix
table(lda.class ,test$Direction)

##
## lda.class Down Up
##      Down   36 47
##      Up     7 14

```

```

# Overall fraction of correct predictions
lda_acc <- mean(lda.class == test$Direction)
lda_acc

## [1] 0.4807692

#-- III) QDA
qda.fit=qda(Direction~Lag1+Lag3+Lag4+Lag5+Volume,
             data=training)
qda.fit

## Call:
## qda(Direction ~ Lag1 + Lag3 + Lag4 + Lag5 + Volume, data = training)
##
## Prior probabilities of groups:
##      Down      Up
## 0.4477157 0.5522843
##
## Group means:
##      Lag1      Lag3      Lag4      Lag5      Volume
## Down  0.28944444 0.17080045 0.15925624 0.21409297 1.266966
## Up    -0.009213235 0.08404044 0.09220956 0.04548897 1.156529

qda.pred=predict(qda.fit,test)
qda.class=qda.pred$class
# Confusion matrix
table(qda.class ,test$Direction)

##
## qda.class Down Up
##      Down  36 53
##      Up    7  8

# Overall fraction of correct predictions
qda_acc <- mean(qda.class == test$Direction)
qda_acc

## [1] 0.4230769

#-- IV) KNN for k=1,2,3,5,7

test.x=cbind(test$Lag1, test$Lag3, test$Lag4, test$Lag5, test$Volume)
training.x=cbind(training$Lag1, training$Lag3, training$Lag4, training$Lag5,
                 training$Volume)
library(class)

k_ls <- c(1,2,3,5,7)
k_acc <- c()
i=0

for(k in k_ls){
  knn.pred=knn(training.x,test.x,training$Direction,k=k)
  contrasts(test$Direction)
  # Confusion matrix
  table(knn.pred ,test$Direction)
  # Overall fraction of correct predictions

```

```

knn_acc <- mean(knn.pred == test$Direction)
k_acc <- c(k_acc,knn_acc)
i=i+1
}

```

```
k_acc
```

```
## [1] 0.4807692 0.5096154 0.5192308 0.4326923 0.4711538
```

```

#Answer:
#Best Variables: Lag1 + Lag3 + Lag4 + Lag5 + Volume
#Best Method: KNN with k=2
#Best confusion matrix:
#knn.pred Down Up
#Down  23 26
#Up    20 35

```

```
### Problem 2 ###
```

```

auto_df <- read.csv('Auto.csv')
#View(auto_df)

```

```

## a)
mpg01 <- rep(0, dim(auto_df)[1])

```

```

# Median of mpg
med_mpg <- median(auto_df$mpg)
med_mpg

```

```
## [1] 22.75
```

```

# mpg01 vector with 0's and 1's
mpg01[auto_df$mpg > med_mpg] = 1
mpg01

```

```

## [1] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 1 1 1 1 1 1 0 0 0 0 0 1 1 1 0 0 0
## [36] 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0
## [71] 0 0 0 0 0 0 0 0 0 1 0 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0
## [106] 0 0 0 0 0 0 0 0 0 1 0 0 1 1 0 0 0 1 0 0 0 0 0 1 1 1 1 0 0 0 0 0 0 0 0 0 1
## [141] 1 1 1 1 1 1 1 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 1 1 1 1 0 1 0
## [176] 1 1 0 1 1 1 1 1 1 1 1 0 0 0 0 0 0 0 1 0 1 1 1 1 0 0 0 0 1 1 1 1 0 0 0 0 0
## [211] 0 0 0 0 1 1 1 1 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1 1 1 0 0 0 1 1 1
## [246] 1 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 0 1 1 1 0 0 0 0 1 1 0 0
## [281] 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0
## [316] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
## [351] 1 1 1 1 1 1 1 1 1 0 1 0 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0 1
## [386] 1 1 1 1 1 1 1

```

```

auto_df <- cbind(auto_df, mpg01)
head(auto_df)

```

```

## mpg cylinders displacement horsepower weight acceleration year origin
## 1 18 8 307 130 3504 12.0 70 1
## 2 15 8 350 165 3693 11.5 70 1
## 3 18 8 318 150 3436 11.0 70 1
## 4 16 8 304 150 3433 12.0 70 1

```

```
## 5 17      8      302      140 3449      10.5 70      1
## 6 15      8      429      198 4341      10.0 70      1
##                                     name mpg01
## 1 chevrolet chevelle malibu      0
## 2      buick skylark 320          0
## 3      plymouth satellite          0
## 4          amc rebel sst          0
## 5          ford torino            0
## 6          ford galaxie 500       0
```

*## b) Explore relationship of mpg01 with other vars*

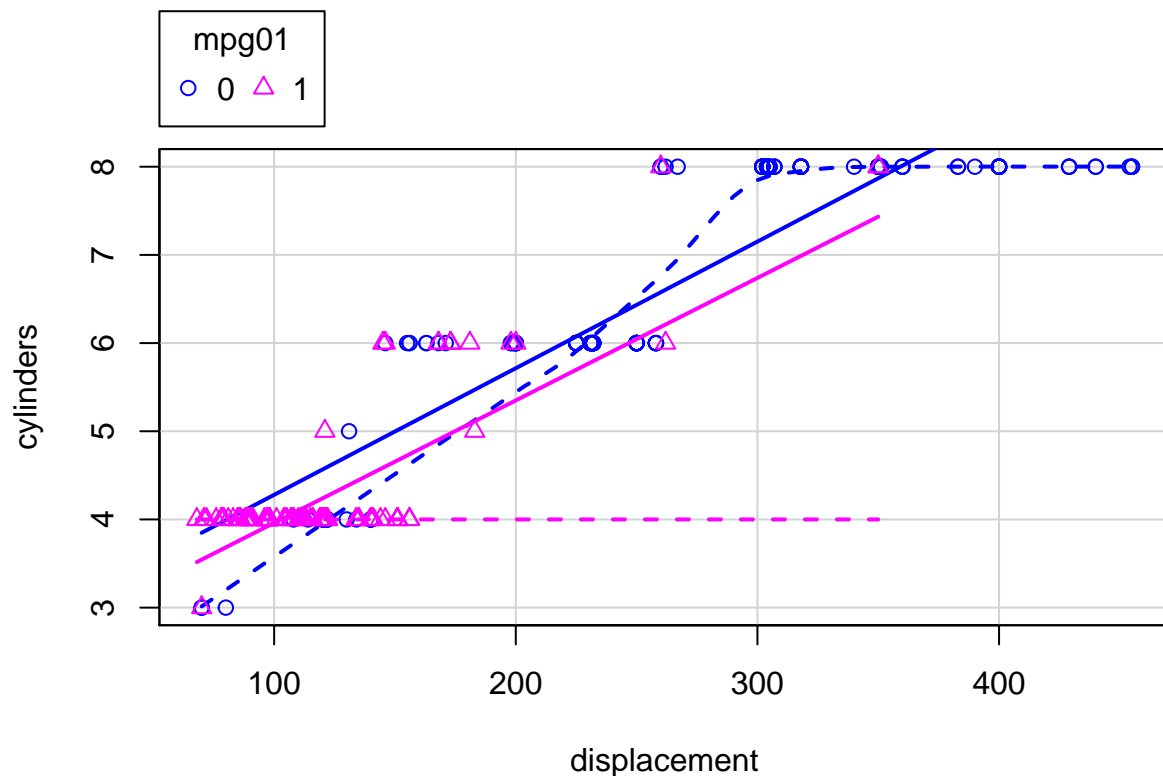
```
library(car)
```

## Loading required package: carData

*## cylinders vs displacement w.r.t mpg01:*

*# We can see that mileage is higher for less number of cylinders and low displacement*  
*# And mileage is lower for higher number of cylinders and higher displacement values*

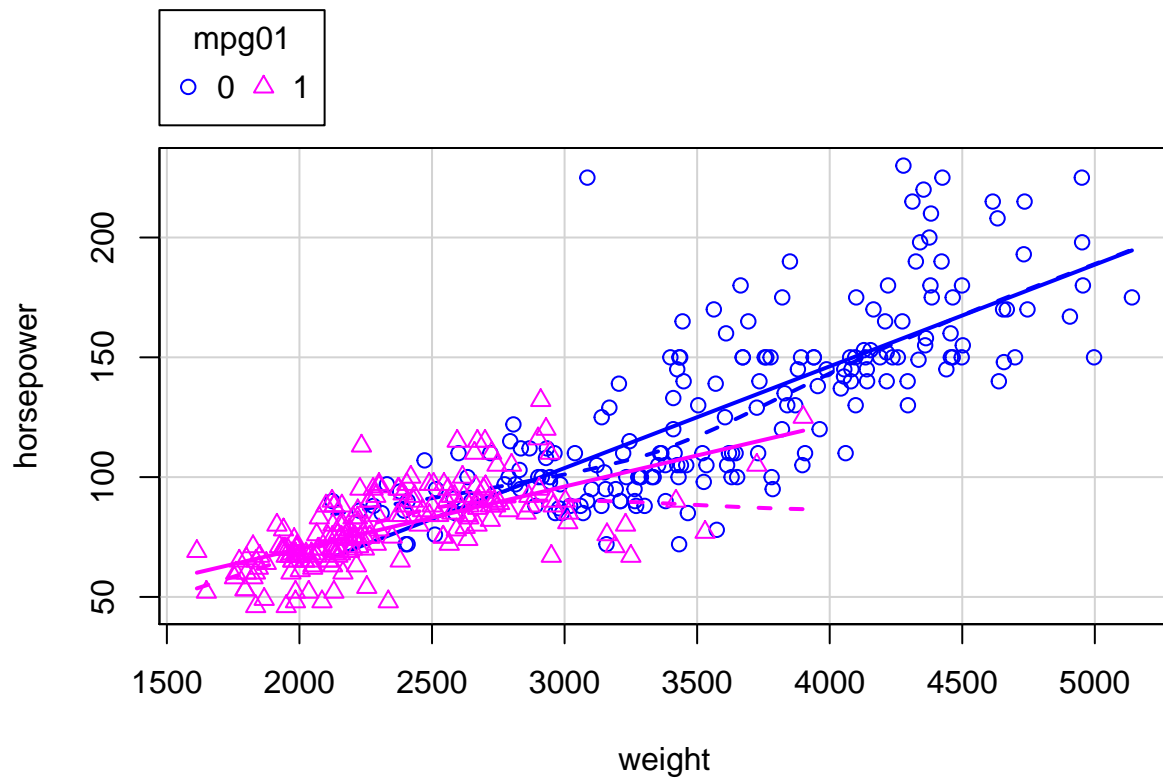
```
scatterplot(cylinders ~ displacement | mpg01, data=auto_df,
            xlab="displacement", ylab="cylinders")
```



*## horsepower vs weight w.r.t mpg01:*


*# From plot: mileage is higher for lower weight values and lower horsepower values*

```
scatterplot(horsepower ~ weight | mpg01, data=auto_df,
            xlab="weight", ylab="horsepower")
```




```
## weight vs acceleration w.r.t. mpg01:  
# From plot: mileage is slightly skewed  
scatterplot(weight ~ acceleration | mpg01, data=auto_df,  
            xlab="acceleration", ylab="weight")
```







```
## year vs origin w.r.t mpg01:  
# From plot: mileage doesn't seem to depend on year or origin, so we can ignore  
# year and origin from predictors  
scatterplot(year ~ origin | mpg01, data=auto_df, xlab="origin", ylab="year")
```




```
# From plot: Lower number cylinders give higher mileage
boxplot(cylinders ~ mpg01, data=auto_df, xlab="mpg01", ylab="cylinders")
```




```
# From plot: Higher displacement implies lower mileage  
boxplot(displacement ~ mpg01, data=auto_df, xlab="mpg01", ylab="displacement")
```




```
# From plot: Mileage is higher for horsepower less than 100  
boxplot(horsepower ~ mpg01, data=auto_df, xlab="mpg01", ylab="horsepower")
```




```
# From plot: Mileage is higher for low weight vehicles  
boxplot(weight ~ mpg01, data=auto_df, xlab="mpg01", ylab="weight")
```



```
# From plot: Mileage doesn't seem to depend much on acceleration  
boxplot(acceleration ~ mpg01, data=auto_df, xlab="mpg01", ylab="acceleration")
```




```
# From plot, we can see that old cars had lower mileage than newer cars  
boxplot(year ~ mpg01, data=auto_df, xlab="mpg01", ylab="year")
```



```
# From plot: Origin doesn't seem to affect mileage much, we can ignore this variable  
# from predictors  
boxplot(origin ~ mpg01, data=auto_df, xlab="mpg01", ylab="origin")
```





*# From the above plots we can see that variables 'origin' and 'year'*  
*# don't have relationship with target 'mpg01'. Hence, we can take rest of the*  
*# variables as predictors for this dataset*

*## c) Train test split*

*# Processing and Scaling*

```
col_names <- c('cylinders', 'displacement', 'horsepower', 'weight',
               'acceleration')
```

```
auto_df_subset <- auto_df[col_names]
```

```
#View(auto_df_subset)
```

*# Normalized subset*

```
auto_df_subset_scaled <- scale(auto_df_subset)
```

```
#View(auto_df_subset_scaled)
```

*# add mpg01 col*

```
auto_df_subset_scaled_final <- cbind(auto_df_subset_scaled, mpg01)
```

```
#View(auto_df_subset_scaled_final)
```

*# Split*

```
split_size = round(dim(auto_df)[1] * 0.8)
```

```
split_size
```

```
## [1] 314
```

```

training=data.frame(auto_df_subset_scaled_final[1:split_size,])
test=data.frame(auto_df_subset_scaled_final[315:392,])

#View(training)
#View(test)
dim(training)

## [1] 314    6

dim(test)

## [1] 78    6

## d) LDA

lda.fit=lda(mpg01~cylinders+displacement+horsepower+weight+acceleration,
            data=training)
lda.fit

## Call:
## lda(mpg01 ~ cylinders + displacement + horsepower + weight +
##      acceleration, data = training)
##
## Prior probabilities of groups:
##          0          1
## 0.6082803 0.3917197
##
## Group means:
##      cylinders displacement horsepower      weight acceleration
## 0  0.7699698    0.7651627    0.688721  0.7687844    -0.365302
## 1 -0.7818856   -0.7861089   -0.671833 -0.8180198     0.332754
##
## Coefficients of linear discriminants:
##              LD1
## cylinders    -0.71689917
## displacement -0.08463165
## horsepower    0.19728376
## weight       -0.93441398
## acceleration -0.04053379

lda.pred=predict(lda.fit,test)
lda.class=lda.pred$class
# Confusion matrix
table(lda.class ,test$mpg01)

##
## lda.class  0  1
##           0  5 11
##           1  0 62

# Overall fraction of correct predictions
lda_acc <- mean(lda.class == test$mpg01)
lda_acc

## [1] 0.8589744

```

```

# Test error for LDA
1-lda_acc

## [1] 0.1410256

## e) QDA

qda.fit=qda(mpg01~cylinders+displacement+horsepower+weight+acceleration,
             data=training)
qda.fit

## Call:
## qda(mpg01 ~ cylinders + displacement + horsepower + weight +
##      acceleration, data = training)
##
## Prior probabilities of groups:
##      0      1
## 0.6082803 0.3917197
##
## Group means:
##      cylinders displacement horsepower      weight acceleration
## 0  0.7699698    0.7651627    0.688721  0.7687844    -0.365302
## 1 -0.7818856   -0.7861089   -0.671833 -0.8180198     0.332754

qda.pred=predict(qda.fit,test)
qda.class=qda.pred$class
# Confusion matrix
table(qda.class ,test$mpg01)

##
## qda.class  0  1
##           0  5 11
##           1  0 62

# Overall fraction of correct predictions
qda_acc <- mean(qda.class == test$mpg01)
qda_acc

## [1] 0.8589744

# Test error for QDA
1-qda_acc

## [1] 0.1410256

## f) Logistic Regression

glm.fit=glm(mpg01~cylinders+displacement+horsepower+weight+acceleration,
            data=training,family=binomial)
summary(glm.fit)

##
## Call:
## glm(formula = mpg01 ~ cylinders + displacement + horsepower +
##      weight + acceleration, family = binomial, data = training)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max

```

```

## -2.3245 -0.2535 -0.0134 0.2985 3.6850
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) -1.6739      0.3362 -4.979 6.39e-07 ***
## cylinders    -0.2858      0.7886 -0.362 0.7171
## displacement -1.1261      1.1715 -0.961 0.3365
## horsepower   -2.2804      0.9551 -2.388 0.0170 *
## weight        -1.4998      0.8714 -1.721 0.0852 .
## acceleration -0.2551      0.3891 -0.656 0.5121
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 420.45  on 313  degrees of freedom
## Residual deviance: 151.72  on 308  degrees of freedom
## AIC: 163.72
##
## Number of Fisher Scoring iterations: 7

glm.probs=predict(glm.fit,newdata = test, type="response")
glm.pred=rep(0,dim(test)[1])
glm.pred[glm.probs>0.5]=1

# Confusion matrix
table(glm.pred,test$mpg01)

##
## glm.pred  0  1
##           0  5 11
##           1  0 62

# Overall fraction of correct predictions
logistic_acc <- mean(glm.pred == test$mpg01)
logistic_acc

## [1] 0.8589744

# test error rate
1-logistic_acc

## [1] 0.1410256

## g) KNN with different k-values: 1, 2, 3, 5, 7
test.x=cbind(test$cylinders, test$displacement, test$horsepower,
             test$weight, test$acceleration)
training.x=cbind(training$cylinders, training$displacement, training$horsepower,
                 training$weight, training$acceleration)
library(class)

# i) k=1
knn.pred=knn(training.x,test.x,training$mpg01,k=1)

# Confusion matrix
table(knn.pred ,test$mpg01)

```

```
##
## knn.pred  0  1
##          0  5 16
##          1  0 57

# Overall fraction of correct predictions
knn_acc <- mean(knn.pred == test$mpg01)
knn_acc

## [1] 0.7948718

# Test error rate
k_1_error <- 1-knn_acc
k_1_error

## [1] 0.2051282

# ii) k=2

knn.pred=knn(training.x,test.x,training$mpg01,k=2)

# Confusion matrix
table(knn.pred ,test$mpg01)

##
## knn.pred  0  1
##          0  5 15
##          1  0 58

# Overall fraction of correct predictions
knn_acc <- mean(knn.pred == test$mpg01)
knn_acc

## [1] 0.8076923

# Test error rate
k_2_error <- 1-knn_acc
k_2_error

## [1] 0.1923077

# iii) k=3

knn.pred=knn(training.x,test.x,training$mpg01,k=3)

# Confusion matrix
table(knn.pred ,test$mpg01)

##
## knn.pred  0  1
##          0  5 14
##          1  0 59

# Overall fraction of correct predictions
knn_acc <- mean(knn.pred == test$mpg01)
knn_acc

## [1] 0.8205128
```

```

# Test error rate
k_3_error <- 1-knn_acc
k_3_error

## [1] 0.1794872

# iv) k=5

knn.pred=knn(training.x,test.x,training$mpg01,k=5)

# Confusion matrix
table(knn.pred ,test$mpg01)

##
## knn.pred  0  1
##          0  5 13
##          1  0 60

# Overall fraction of correct predictions
knn_acc <- mean(knn.pred == test$mpg01)
knn_acc

## [1] 0.8333333

# Test error rate
k_5_error <- 1-knn_acc
k_5_error

## [1] 0.1666667

# v) k=7

knn.pred=knn(training.x,test.x,training$mpg01,k=7)

# Confusion matrix
table(knn.pred ,test$mpg01)

##
## knn.pred  0  1
##          0  5 13
##          1  0 60

# Overall fraction of correct predictions
knn_acc <- mean(knn.pred == test$mpg01)
knn_acc

## [1] 0.8333333

# Test error rate
k_7_error <- 1-knn_acc
k_7_error

## [1] 0.1666667

# *) comparing all k values
cbind(k_1_error, k_2_error, k_3_error, k_5_error, k_7_error)

##      k_1_error k_2_error k_3_error k_5_error k_7_error
## [1,] 0.2051282 0.1923077 0.1794872 0.1666667 0.1666667

```

*# As we can see,  $k=2$  gives the lowest error rate*